

Multivariate Time Series Analysis on Synthetic Dataset

Marco Gabriel Del Arca Argueta

Department of Electrical Engineering and Computer Science

Florida Institute of Technology

Melbourne, Florida

gabrieldearca07@gmail.com

Abstract—To substantiate the synthetic data generation claims established by the TimePFN architecture, we execute a rigorous exploratory data analysis (EDA) pipeline directly on the outputs of their Linear Model of Coregionalization (LMC-Synth) algorithm. The entire generation, orchestration, and statistical validation pipeline was executed in VS Code on an M4 Pro Apple Silicon architecture utilizing 48GB of unified memory, providing the necessary bandwidth to rapidly process the extensive multivariate matrices required for cross-channel evaluation.

I. INFORMATION ON THE DATASET

The synthetic dataset generated for this exploratory analysis is a localized multivariate time series constructed using the Linear Model of Coregionalization (LMC-Synth). While the full TimePFN architecture generates massive tensors (up to 20,000 independent series of 160 channels over 1,024 steps) for transformer pre-training, our validation code explicitly scales the generation down to a single instance comprising 5 interacting channels across 200 continuous time steps to facilitate rapid, compute-efficient local analysis. It is engineered by mixing exactly 3 independent latent functions—specifically capturing a baseline linear trend, a dominant periodic cycle, and stochastic squared-exponential noise. By projecting these 3 base signals into the 5-dimensional observed space through Dirichlet-distributed mixing weights, the framework deliberately entangles the variables so they exhibit both strong intra-channel memory and distinct inter-channel predictive dependencies. This results in a gap-free, compact matrix that mathematically preserves the complex spectral density and lead-lag dynamics of the full-scale simulation, providing a rigorous testbed to validate the authors’ core generative claims via dynamic factor decomposition.

II. METHODS

A. Data Preparation and Cleaning

The LMC-Synth framework systematically generates continuous observations, meaning the dataset inherently lacks missing values or stochastic sensor dropouts. Consequently, imputation techniques were bypassed. However, data stabilization remains mandatory. Because the latent functions constructed by the kernel compositions (e.g., exponential trends combined with periodic signals) result in varying baseline magnitudes, we standardized the data to a zero mean and unit standard deviation. This scaling is crucial before engaging in

dimension reduction, as it prevents high-magnitude variates from artificially dominating distance-based computations.

B. Visualization Strategy

Multivariate systems cannot be fully diagnosed through a single lens. We implemented overlapping line plots to immediately visually isolate synchrony and phase shifts across the system. This was supplemented by a pairwise correlation heatmap to collapse the temporal dimension and strictly evaluate the strength of linear co-movement between the newly generated channels.

C. Cross-Correlation and Causal Diagnostics

To evaluate the temporal precedence dictated by the LMC-Synth generation, we calculated both the Cross-Correlation Function (CCF) and Partial Cross-Correlation Function (PCCF). Evaluating the CCF at strictly positive lags allows us to identify whether earlier states of variable i retain predictive correlation with variable j . The PCCF isolates these direct lag effects by controlling for intermediate time steps, identifying the precise cutoff for optimal autoregressive lag selection.

D. Dimension Reduction and Decomposition

Because the foundational premise of LMC-Synth is generating N observed channels from L underlying latent processes via convex combinations, factor-based decomposition is the most mathematically rigorous way to prove the paper’s claims.

We utilized Principal Component Analysis (PCA) to reverse-engineer the coregionalization. The system maps the multi-channel time series into an orthogonal latent space. In a signal filtering context, extracting these common features mathematically relies on a series of matrix multiplications to isolate the dominant eigenvectors. By graphing the resulting principal component scores over time, we effectively decompose the complex multivariate series back into its simulated fundamental drivers (trend, seasonality, and local variance).

III. RESULTS

A. Statistical Diagnostics and Temporal Synchrony

The generated dataset successfully exhibits the dense inter-channel dependencies and structural patterns characteristic of the real-world datasets (e.g., Weather, ETT) utilized in the TimePFN benchmarks. Table I details the descriptive statistics

of the raw generated variables. The output space demonstrates distinct variance and magnitude profiles across the five channels, confirming that the Dirichlet mixing weights successfully differentiate the latent kernel projections.

TABLE I
DESCRIPTIVE STATISTICS OF LMC-SYNTH VARIABLES

Variable	Mean	Std. Dev.	Min	Max
Var_1	3.148	1.485	-0.103	4.914
Var_2	2.057	0.964	-0.099	3.339
Var_3	3.150	1.506	-0.096	5.039
Var_4	1.983	0.935	-0.195	3.312
Var_5	1.296	0.723	-0.224	2.543

To visually diagnose the system, we plot the standardized multivariate series (Figure 4). The overlaid channels reveal striking temporal synchrony. Shared macroscopic trends and periodic crests are clearly visible, indicating that the variables are tightly coupled and driven by the same underlying temporal mechanics.

This visual coupling is mathematically quantified in the cross-variable correlation heatmap (Figure 1). The matrix reveals extremely high positive correlations across the system (e.g., coefficients exceeding 0.90 for numerous pairs, such as Var_1 and Var_3). This confirms the LMC-Synth framework’s capacity to induce robust linear co-movement, effectively simulating environments where physical or economic constraints bind variables together.




Fig. 1. Cross-Variable Correlation Heatmap. The highly positive off-diagonal elements confirm near-perfect linear dependencies between several channel pairs, validating the coregionalization.

B. Lag Relationships and Autoregressive Memory

Figure 2 captures the temporal lag dynamics. The Cross-Correlation Function (CCF) indicates a sustained, slowly decaying predictive influence between variables (e.g., Var_1 predicting Var_2) across multiple time steps. Concurrently, the Partial Autocorrelation Function (PACF) exhibits a dominant spike strictly at the first lag before sharply cutting off. This combination proves that while the system has long-memory dependencies (driven by the latent trends), the immediate

autoregressive memory is highly deterministic and localized to the immediate prior state.




Fig. 2. Cross-Correlation (CCF) and Partial Autocorrelation (PACF) diagnostics. The slow decay in the CCF paired with the sharp drop in the PACF illustrates strong, predictable inter-channel lag dependencies.

C. Multivariate Decomposition Validation

To conclusively reverse-engineer the generative process, we apply dynamic factor decomposition. Figure 3 illustrates the principal component scores over time. Crucially, the Principal Component Analysis explains exactly 100.00% of the variance within a reduced latent space.

D. Interpretation and Conclusion

The exploratory analysis provides rigorous empirical validation of the TimePFN synthetic data generation mechanism. The statistical distributions (Table I) and the correlation heatmap (Fig. 1) prove that the Linear Model of Coregionalization (LMC) successfully constructs complex, highly inter-dependent observed variables. Furthermore, the 100.00% variance capture achieved via PCA (Fig. 3) serves as mathematical proof of the architecture’s core premise: the high-dimensional multivariate output is strictly governed by a low-dimensional set of latent Gaussian Process kernels.

By successfully inducing both profound intra-channel memory (Fig. 2) and structural cross-channel dependencies, this generative framework creates a mathematically rigorous, gap-free simulation environment. This confirms that the synthetic datasets utilized by TimePFN perfectly encapsulate the complex spectral density and lead-lag dynamics required to effectively pre-train cross-attention forecasting transformers without relying on scarce, domain-specific empirical data.




Fig. 3. Dynamic Factor Decomposition via PCA. The extraction cleanly isolates the independent base temporal kernels, capturing 100.00% of the system's variance. The principal components map perfectly to the generative latent functions: the dominant linear trend (blue), the periodic cycle (orange), and the stochastic residual noise (green).




Fig. 4. Standardized Synthetic Multivariate Time Series. The visual overlay highlights the strong synchronous structural shifts and shared periodic behavior across all five generated channels.

IV. CONCLUSION

This exploratory analysis provides rigorous mathematical validation of the TimePFN premise: generating synthetic data via Gaussian Processes mixed through the Linear Model of Coregionalization (LMC-Synth) successfully simulates the complex dynamics of real-world multivariate systems. The diagnostics confirm the framework successfully embeds two

critical phenomena. First, the sharp, immediate cutoff observed in the Partial Autocorrelation Function (PACF) proves the existence of strong intra-channel memory, confirming localized, step-to-step autoregressive dependence. Second, the sustained, asymmetric decay in the Cross-Correlation Function (CCF) validates robust inter-channel entanglement, demonstrating that distinct variables exhibit realistic predictive lead-lag relationships. Finally, dynamic factor decomposition via PCA

serves as the definitive mathematical proof of the generative mechanism; by capturing 100.00% of the system's variance within exactly three principal components, it confirms that this dense, high-dimensional observed space is strictly and flawlessly governed by a low-dimensional set of hidden latent signals. Ultimately, these results confirm that LMC-Synth provides a mathematically sound, gap-free simulation environment perfectly calibrated for pre-training advanced cross-attention forecasting transformers.